

# Discussion April 20: Homework 1 Tutorial

## 1. Obtaining and Reading the Data

### 1.1. Obtaining the Data

- **Problem 1 (1).** The data is called *cpusmall*. The data is presented on an html page. You can copy and paste the data into a blank text document and save it as *cpusmall.txt*.
- **Problem 1 (3).** The data are *E2006.train* and *E2006.test* respectively. Start by downloading the .bz2 files. The .bz2 file extension is a compression format mainly used in Linux systems. Go here for a list of all programs in Windows and Mac that open this file extension: <https://fileinfo.com/extension/bz2>. After extracting the file, there is no need to rename them (they can end in the extension .train and .test).
- **Problem 2 (2).** The data is *news20*. The steps are similar to Problem 1 (3). There is also no need to rename this file (it will end in the extension .binary).

### 1.2. Reading the Data into Python

All of the data in this homework are in a specific format called *svmlib*. The easiest way to read the data into Python is using a function in the datasets package from *scikit-learn*. If you have not done so, you will need to install scikit-learn first in order to do this. The following example will read in *E2006.train*.

```
from sklearn import datasets

filename = "C:/Users/jstwa/Desktop/STA 141C/E2006.train"
X,y = datasets.load_svmlight_file(filename)
```

The feature matrix  $X$  is in a weird format called CSR (Compressed Sparse Row). In order to make it easier to work with, we will convert it into a regular numpy 2D array. This process involves using scipy:

```
from scipy import sparse

X_array = sparse.csr_matrix.todense(X)
```

## 2. Machine Learning Concepts Review

### 2.1. Gradient Descent

This is an optimization algorithm used to obtain a local minimum of a given function. It is popular in machine learning and is usually used when an analytical solution is not possible. Here is the algorithm in pseudocode:

```
init  $w_i^{(0)}$  = random value for each  $i = 1, \dots, p$   
 $r_0 = \|\nabla f(\mathbf{w}^{(0)})\|$ ,  $\eta$  = fixed step size chosen in the interval  $(0, 1)$   
for  $t$  in  $\text{range}(0, 200)$  or until  $\|\nabla f(\mathbf{w}^{(t)})\| < \epsilon \cdot r_0$  do:  
     $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \eta \cdot \nabla f(\mathbf{w}^{(t)})$ 
```

Gradient descent is very easy to implement. The main thing you will have to do is to (1) Find an expression for  $\nabla f(w_t)$ , and (2) Implement this expression in Python. I recommend doing (2) as its own function so the code will be cleaner.

### 2.2. Cross-Validation

5-fold Cross-validation is a way to validate your model; i.e. check its prediction power. Here is a rough pseudocode for Problem 1 (2):

Divide your data as evenly as possible into 5 pieces:  $D_0, \dots, D_4$ .

```
for  $i$  in  $\text{range}(0, 4)$  do:  
    test.data =  $D_i$   
    train.data = Everything else except  $D_i$   
     $w.\text{train}$  = ridge.regression(train.X)  
     $\text{mse}_i$  = mse( $x$  = test.X,  $y$ =test.y,  $w$  =  $w.\text{train}$ )  
  
 $\text{mse}_{cv}$  = average( $\text{mse}_i$ ,  $i = 0, \dots, 4$ )
```

### 3. Deriving the Gradient for Ridge Regression

**3.1. Reminder: What is a gradient?** Suppose that  $f$  is a function that takes a  $p \times 1$  vector as an input and returns a single number (scalar). The gradient of  $f$ , denoted by  $\nabla f$ , is given by:

$$\nabla f = \nabla f(w_1, \dots, w_p) = \left[ \frac{\delta f}{\delta w_1}, \dots, \frac{\delta f}{\delta w_p} \right]$$

In other words, it is a vector containing all the partial derivatives for each  $w_i$  in  $f$ .

**3.2. Objective function for ridge regression.** For ridge regression, letting  $\lambda = 1$ , we wish to minimize the following function:

$$f(w_1, \dots, w_p) = \frac{1}{2} \sum_{i=1}^n (x_{i1}w_1 + \dots + x_{ip}w_p - y_i)^2 + \frac{1}{2} (w_1^2 + \dots + w_p^2)$$

Let's examine the partial derivative of the above with respect to an arbitrary  $w_j$ :

$$\frac{\delta f}{\delta w_j} = \sum_{i=1}^n (x_{i1}w_1 + \dots + x_{ip}w_p - y_i) (x_{ij}) + w_j$$

So overall, the gradient looks like this:

$$\begin{aligned} \nabla f &= \left[ \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{w} - y_i) (x_{i1}) + w_1, \dots, \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{w} - y_i) (x_{ip}) + w_p \right] \\ &= \left[ \sum_{i=1}^n \mathbf{x}_i^T \mathbf{w} \cdot x_{i1} - \sum_{i=1}^n y_i x_{i1} + w_1, \dots, \sum_{i=1}^n \mathbf{x}_i^T \mathbf{w} \cdot x_{ip} - \sum_{i=1}^n y_i x_{ip} + w_p \right] \\ &= \left[ \sum_{i=1}^n \mathbf{x}_i^T \mathbf{w} \cdot x_{i1}, \dots, \sum_{i=1}^n \mathbf{x}_i^T \mathbf{w} \cdot x_{ip} \right] - \left[ \sum_{i=1}^n y_i x_{i1}, \dots, \sum_{i=1}^n y_i x_{ip} \right] + [w_1, \dots, w_p] \\ &= \mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{X}^T \mathbf{y} + \mathbf{w} \\ &= \mathbf{X}^T (\mathbf{X} \mathbf{w} - \mathbf{y}) + \mathbf{w} \end{aligned}$$